

ORACLE®

# MySQL 5.7 GIS

Norvald H. Ryeng  
Software Engineer  
MySQL Optimizer Team

Trondheim, March 2015



# Agenda

- 1 What is GIS?
- 2 Basic concepts
- 3 GIS functions
- 4 Indexing and optimization
- 5 Displaying data in a web map
- 6 The future of MySQL GIS





# What is GIS?

“The early days of GIS were very lonely.  
No-one knew what it meant.”

— Roger Tomlinson, “Father of GIS”

# What is GIS?

Geocoding  
Satellite imagery  
Decision analysis **Routing** **Raster data**  
2d Network models 4d  
Vector data **Navigation** **Web maps** **Monitoring**  
**Cartography** Topology Hydrology  
Statistics Positioning **3d** **City planning**  
Remote sensing  
Standards **Maps** Modelling **Projections**  
**Overlays**

“A geographic information system (GIS) is a system designed to capture, store manipulate, analyze, manage and present all types of spatial or geographical data.”

— Wikipedia

# What is GIS?

A word cloud featuring various GIS software and tools. The words are arranged in a non-uniform, overlapping manner, with some appearing larger and bolder than others. The colors are primarily shades of gray, with some words in red. The words include:

- FalconView
- OGR
- SuperMap
- Mapnik
- Kalypso
- Google Earth
- uDig
- Bing Maps
- Smallworld
- ILWIS
- Galileo
- OpenLayers
- MySQL Workbench
- MapFish
- GeoServer
- TerraView
- MapServer
- OpenStreetMap
- GPS
- GDAL
- Google Maps
- MySQL
- ArcGIS
- AutoCAD
- GRASS
- SAGA
- NetCAD



# GIS in database management systems

- Geometric shapes
  - Points, lines, polygons
- Functions
  - Comparisons
  - Generate new shapes
  - Measures
  - Properties
- Data types
- Indexes
- Query optimization
- Standards
  - SQL/MM
  - OGC Simple Feature Access
- Coordinate systems

A man and a woman are working together at a wooden table. The man, wearing a blue and white striped cardigan, is seated and looking down at a set of architectural plans. The woman, wearing a yellow sweater, is standing and leaning over the table, holding a pen and looking at the plans. On the table, there is a yellow cup, a rolled-up blueprint, and a ruler. The background shows a large window with a view of a city street.

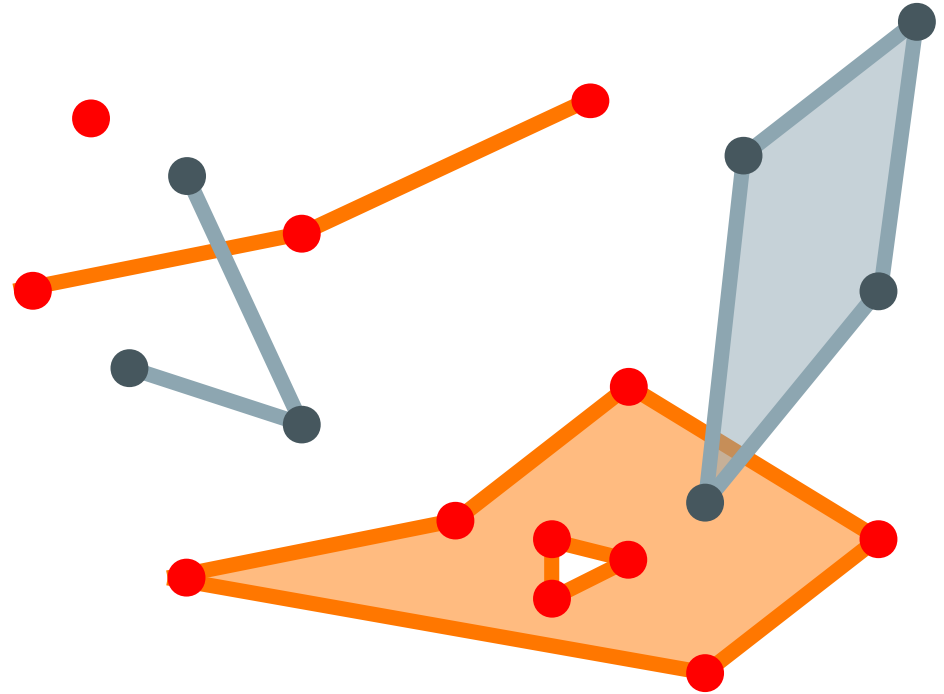
# Basic concepts

“Geography is just physics slowed down,  
with a couple of trees stuck in it.”

— Terry Pratchett, in *The Last Continent*

# Geometric objects

- Point
- LineString
- Polygon
- MultiPoint
- MultiLineString
- MultiPolygon
- GeometryCollection



# Point



Point(0, 0)

```
ST_GeomFromText('POINT(0 0)')
```

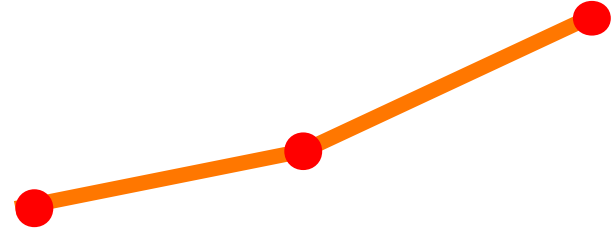
```
ST_GeomFromGeoJSON('{"type":"Point","coordinates":[0,0]}')
```

```
ST_PointFromGeohash('s000', 0)
```

ST\_GeomFromWKB(0x010100000000000000000000000000000000000000)

# LineString

- Two or more points



```
LineString(Point(0, 0), Point(1, 1))
```

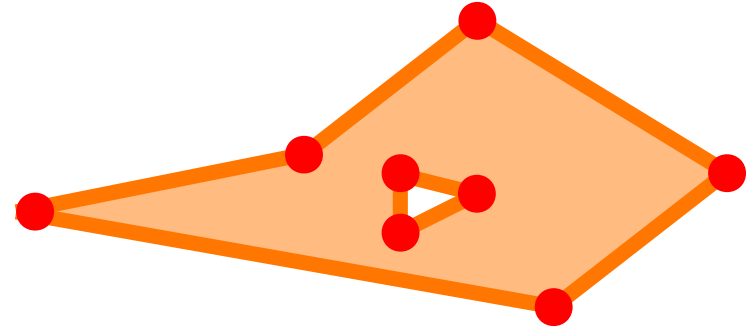
```
ST_GeomFromText('LINESTRING(0 0, 1 1)')
```

```
ST_GeomFromGeoJSON('{"type":"LineString","coordinates":[[0,0],[1,1]]}')
```



# Polygon

- One exterior ring
- Zero or more inner rings (holes)
- At least four points in each ring
  - Start and end point is the same



Inner ring

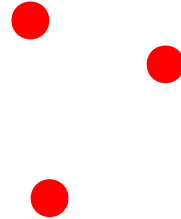
```
Polygon(LineString(Point(0, 0), Point(1, 0), Point(1, 1), Point(0, 0)))
```

```
ST_GeomFromText('POLYGON((0 0, 1 0, 1 1, 0 0), (0.2 0.1, 0.9 0.8, 0.9 0.1, 0.2 0.1))')
```

```
ST_GeomFromGeoJSON('{"type":"Polygon","coordinates":[[[0,0],[1,0],[1,1],[0,0]]]}')
```

# MultiPoint

- One or more Point
  - Can't be empty



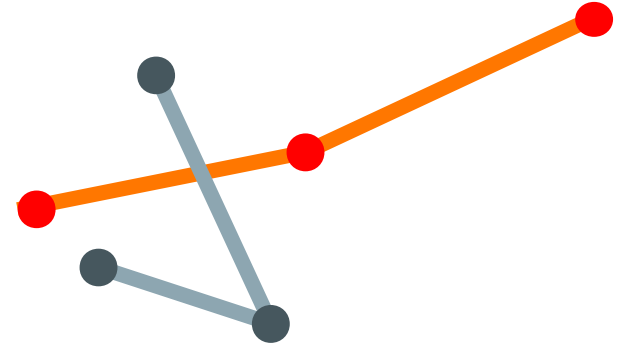
```
MultiPoint(Point(0, 0), Point(1, 1))
```

```
ST_GeomFromText('MULTIPOINT(0 0, 1 1)')
```

```
ST_GeomFromGeoJSON('{"type":"MultiPoint","coordinates":[[0,0],[1,1]]}')
```

# MultiLineString

- One or more LineString
  - Can't be empty



```
MultiLineString(LineString(Point(0, 0), Point(1, 1)), LineString(Point(2, 2), Point(3, 3)))
```

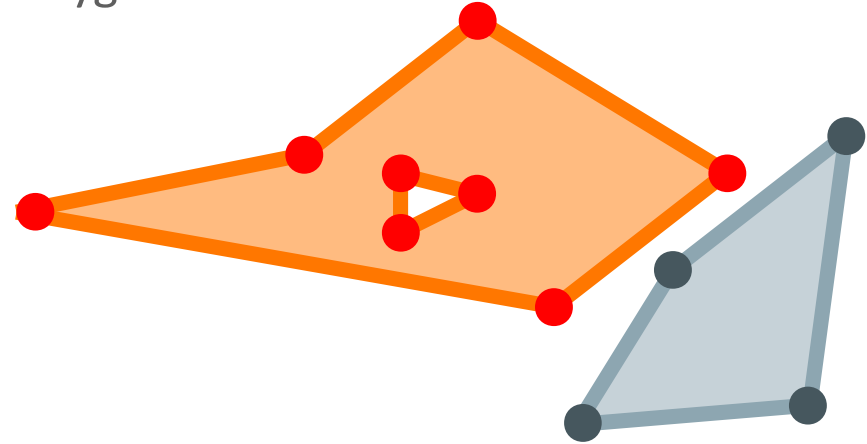
```
ST_GeomFromText('MULTILINESTRING((0 0, 1 1),(2 2, 3 3))')
```

```
ST_GeomFromGeoJSON('{"type":"MultiLineString","coordinates":[[[0,0],[1,1]],[[2,2],[3,3]]]}')
```

# MultiPolygon

Will make the MultiPolygon invalid

- One or more Polygon
  - Can't be empty
- Polygons shouldn't overlap
  - May touch in a finite number of points



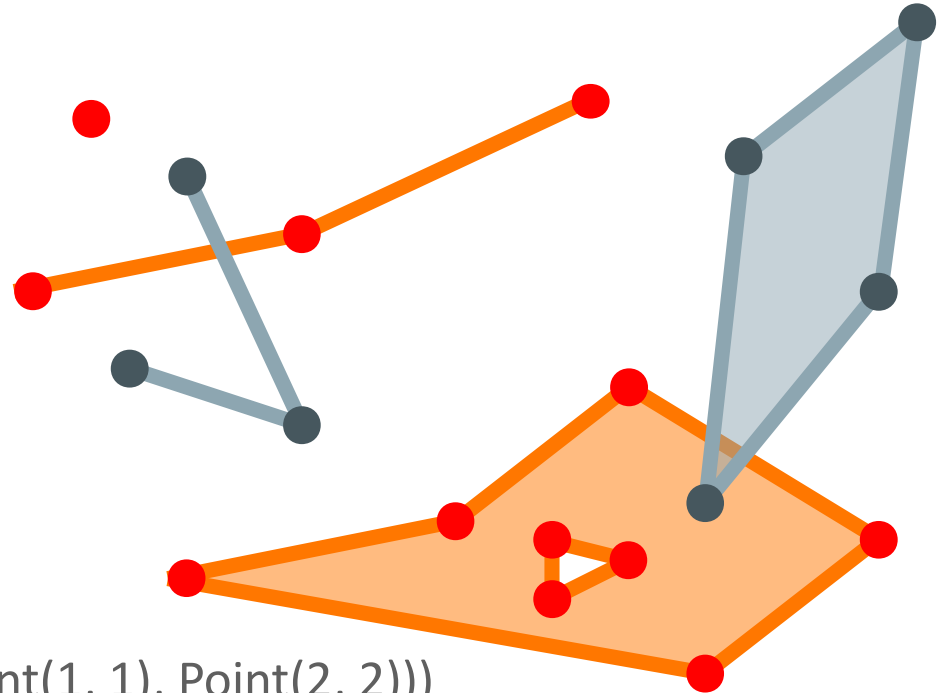
```
MultiPolygon(Polygon(LineString(Point(0, 0), Point(1, 0), Point(1, 1), Point(0, 0))))
```

```
ST_GeomFromText('MULTIPOLYGON(((0 0, 1 0, 1 1, 0 0)))')
```

```
ST_GeomFromGeoJSON('{"type":"MultiPolygon","coordinates":[[[[[0,0],[1,0],[1,1],[0,0]]]]}')
```

# GeometryCollection

- Zero or more geometries
  - May be empty
- No restrictions on overlapping



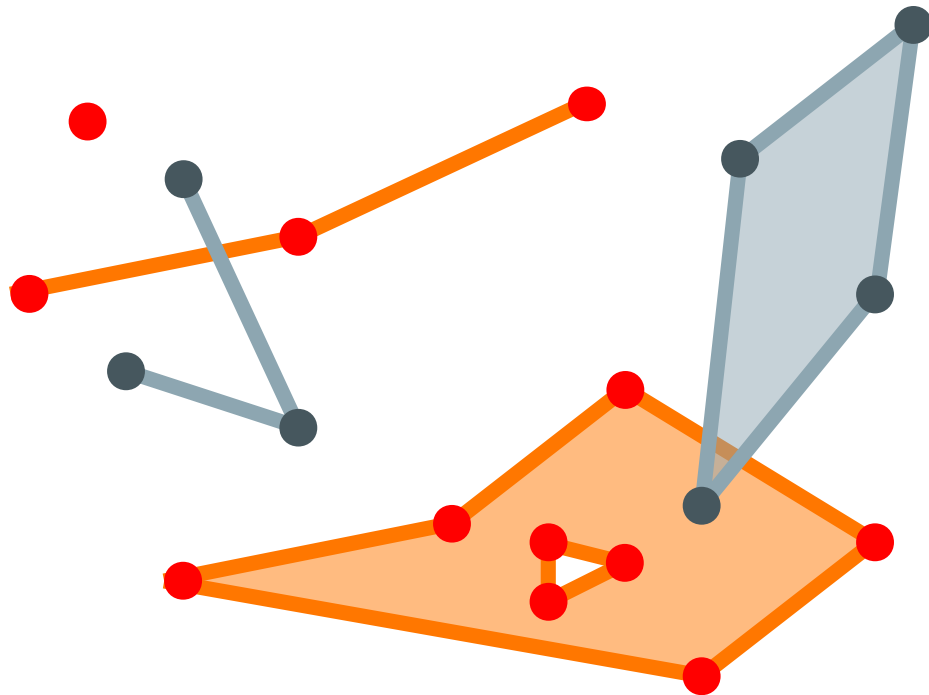
```
GeometryCollection(Point(0, 0), Linestring(Point(1, 1), Point(2, 2)))
```

```
ST_GeomFromText('GEOMETRYCOLLECTION(POINT(0 0), LINESTRING(1 1, 2 2))')
```

```
ST_GeomFromGeoJSON('{"type":"GeometryCollection","geometries":  
[{"type":"Point","coordinates":[0,0]},{"type":"LineString","coordinates":[[1,1],[2,2]]}]}')
```

# Data types

- Geometry (may store any of the types below)
- Point
- LineString
- Polygon
- MultiPoint
- MultiLineString
- MultiPolygon
- GeometryCollection





# Example

- Sightseeing in Trondheim
- Database of popular places to visit
  - A unique ID
  - Position (point)
  - Descriptive text

# Example

```
CREATE TABLE sights (  
  id INT AUTO_INCREMENT PRIMARY KEY,  
  pos POINT NOT NULL,  
  description VARCHAR(200)  
) ENGINE=InnoDB;
```

```
INSERT INTO sights (pos, description) VALUES (  
  Point(10.3958, 63.4269), 'Nidaros Cathedral'  
);
```

```
SELECT ST_AsText(pos), description FROM sights;
```

<b>ST_AsText(pos)</b>	<b>description</b>
POINT(10.3958 63.4269)	Nidaros Cathedral

# Coordinate systems

- Each geometry is in a spatial reference system (SRS)
  - Specified by SRID (integer)
  - Geometries in different SRSs can't be compared
- MySQL supports a Cartesian plane (SRID 0)
  - Default if no SRS is specified
    - Unless the import format defaults to another SRS, e.g., WGS 84 for GeoJSON
- Other reference systems
  - Typically defined by EPSG, e.g., WGS 84 (SRID 4326)
  - Computations are always done in SRID 0

# Longitude and latitude

- OGC specifies X and Y axes, not longitude and latitude
- In GIS,  $\langle \text{longitude}, \text{latitude} \rangle$  is the de facto standard
  - X = degrees East (negative for West)
  - Y = degrees North (negative for South)

# Example

```
INSERT INTO sights (pos, description) VALUES (  
    ST_GeomFromGeoJSON('{"type":"Point","coordinates":[10.4025,63.4194]}'),  
    'Norwegian University of Science and Technology'  
);
```

```
INSERT INTO sights (pos, description) VALUES (  
    ST_GeomFromText('POINT(10.3948 63.4225)', 4326),  
    'Student Society Building'  
);
```

```
INSERT INTO sights (pos, description) VALUES (  
    ST_GeomFromText('POINT(10.3951 63.4305)'),  
    'Olav Tryggvason Monument'  
);
```

# Example

```
SELECT ST_AsText(pos), ST_SRID(pos), description FROM sights;
```

<b>ST_AsText(pos)</b>	<b>ST_SRID(pos)</b>	<b>description</b>
POINT(10.3958 63.4269)	0	Nidaros Cathedral
POINT(10.4025 63.4194)	4326	Norwegian University of Science and Tech ...
POINT(10.3969 63.428)	4326	Student Society Building
POINT(10.3951 63.4305)	0	Olav Tryggvason Monument



# GIS functions

A large concrete bridge pier is shown under construction. The pier has a wide, flat top surface and a thick, tapered base. Scaffolding and safety railings are visible on the top surface. In the background, a white crane stands next to another bridge pier. The sky is blue with scattered white clouds. The overall scene is an industrial construction site.

# Functions

- SQL/MM style naming
  - ST\_ prefix
- MBR functions
  - Minimum bounding rectangle instead of exact shape
- Geometry construction functions
  - Same name as data type

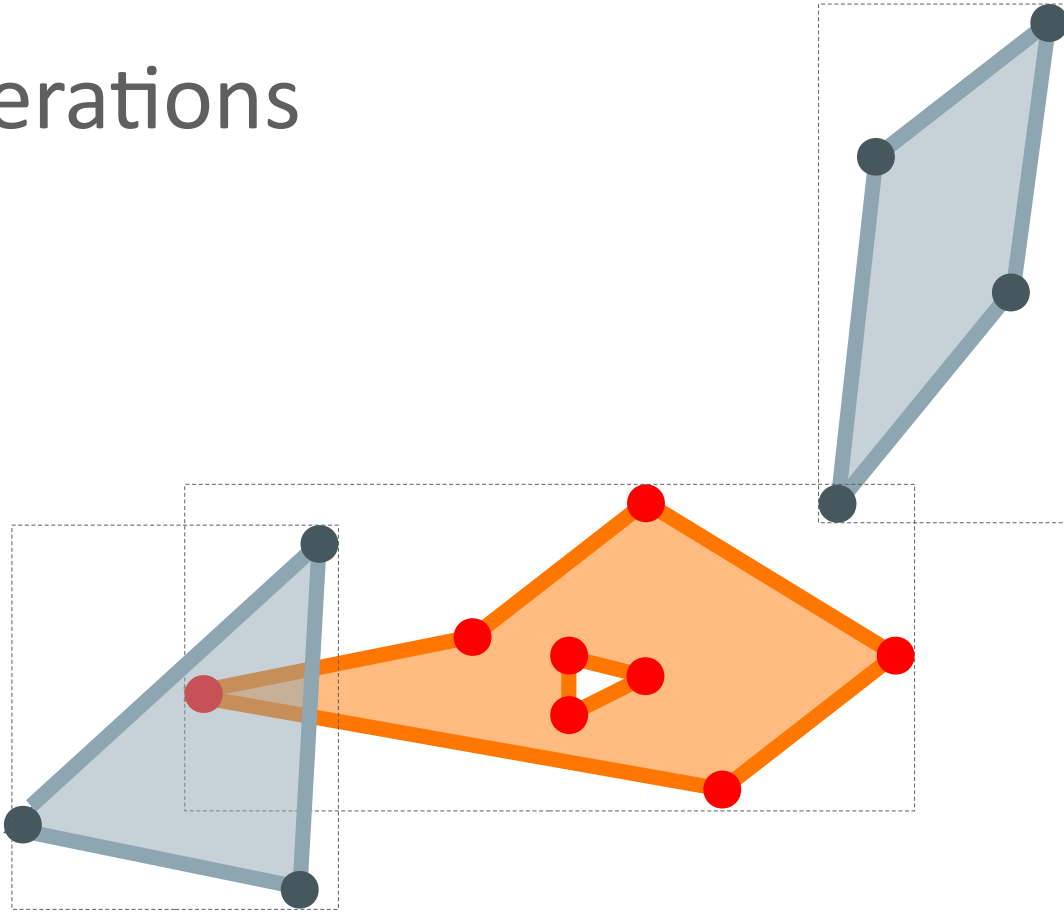
# Functions

- Only defined for valid geometries
- The result is undefined for invalid geometries
  - A best effort result
  - A weird result
  - An error
- Check with ST\_IsValid or use ST\_Validate if unsure
- Mixing SRIDs is not allowed
  - Results in an error

# Functions

- Comparison
  - ST\_Contains, ST\_Crosses, ST\_Disjoint, ST\_Equals, ST\_Intersects, ST\_Overlaps, ST\_Touches, ST\_Within
- Produce new geometries
  - ST\_Buffer, ST\_Centroid, ST\_ConvexHull, ST\_Envelope, ST\_MakeEnvelope, ST\_Simplify, ST\_Difference, ST\_Intersection, ST\_SymDifference, ST\_Union

# Set operations



# Functions

- Measures
  - ST\_Area, ST\_Distance, ST\_Distance\_Sphere, ST\_Length
- Extract properties
  - ST\_Dimension, ST\_EndPoint, ST\_ExteriorRing, ST\_GeometryN, ST\_GeometryType, ST\_InteriorRingN, ST\_IsClosed, ST\_IsEmpty, ST\_IsSimple, ST\_IsValid, ST\_PointN, ST\_SRID, ST\_StartPoint, ST\_X, ST\_Y
- Helper functions
  - ST\_LatFromGeohash, ST\_LongFromGeohash, ST\_Validate



# Functions

- Import

- ST\_GeomCollFromTxt/ST\_GeomCollFromText, ST\_GeomCollFromWKB, ST\_GeomFromGeoJSON, ST\_GeomFromText, ST\_GeomFromWKB, ST\_LineFromText, ST\_LineFromWKB, ST\_MLineFromText, ST\_MLineFromWKB, ST\_MPointFromText, ST\_MPointFromWKB, ST\_MPolyFromText, ST\_MPolyFromWKB, ST\_PointFromGeohash, ST\_PolyFromText, ST\_PolyFromWKB

- Export

- ST\_AsBinary, ST\_AsGeoJSON, ST\_AsText, ST\_Geohash

# Example

```
SET @city_center = ST_GeomFromText(  
    'POLYGON((10.3765 63.4292, 10.3847 63.4277, 10.3902 63.4247, 10.3986 63.4245,  
    10.4013 63.4264, 10.4013 63.4283, 10.4072 63.4347, 10.4037 63.4354,  
    10.3954 63.4350, 10.3799 63.4314, 10.3765 63.4292))'  
);
```

```
SELECT description FROM sights  
WHERE ST_Within(pos, @city_center);
```

# Example

```
SET @city_center = ST_GeomFromText(  
    'POLYGON((10.3765 63.4292, 10.3847 63.4277, 10.3902 63.4247, 10.3986 63.4245,  
    10.4013 63.4264, 10.4013 63.4283, 10.4072 63.4347, 10.4037 63.4354,  
    10.3954 63.4350, 10.3799 63.4314, 10.3765 63.4292))'  
);
```

```
SELECT description FROM sights  
    WHERE ST_Within(pos, @city_center);
```

**ERROR 3033 (HY000): Binary geometry function st\_within given two geometries of different srids: 4326 and 0, which should have been identical.**

# Example

```
UPDATE sights SET pos = ST_GeomFromWKB(ST_AsBinary(pos));
```

```
SELECT description FROM sights  
  WHERE ST_Within(pos, @city_center);
```

## **description**

Nidaros Cathedral

Olav Tryggvason Monument

# Example

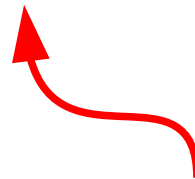
Standing at the Olav Tryggvason Monument, which sights are nearby?

```
SELECT description FROM sights  
  WHERE ST_Within(pos, ST_Buffer(Point(10.3951, 63.4305), 0.004));
```

## **description**

Nidaros Cathedral

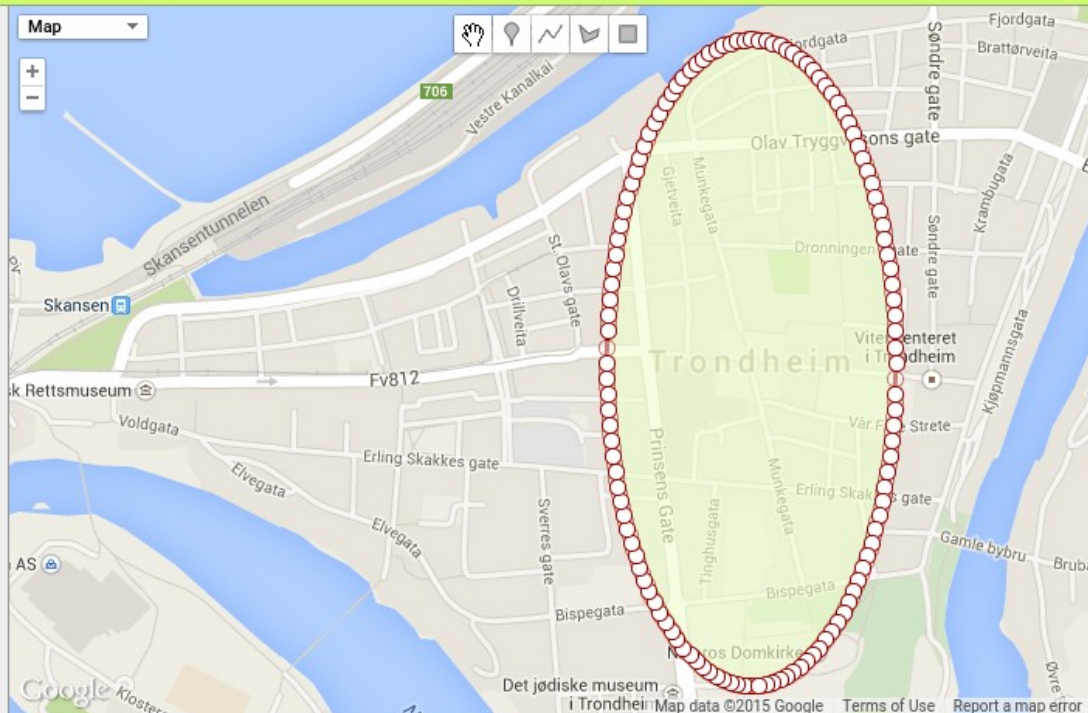
Olav Tryggvason Monument



“nearby” = within 0.004 degrees

# Wicket

It whispers WKT in your application's ear.



Wicket is a lightweight Javascript library that reads and writes **Well-Known Text (WKT)** strings. It can also be extended to parse and to create geometric objects from various mapping frameworks, such as **Leaflet** and the Google Maps API.

```
9012
63.426619874987225,10.39588036128
8064
63.42657685887839,10.395686921897
822
63.426543293960144,10.39549206856
1317
63.42651926109331,10.395296270697
308 63.42650481817518,10.3951
63.426500000000004))
```

☐ Format for URLs

Clear Map

Map It!

# Example

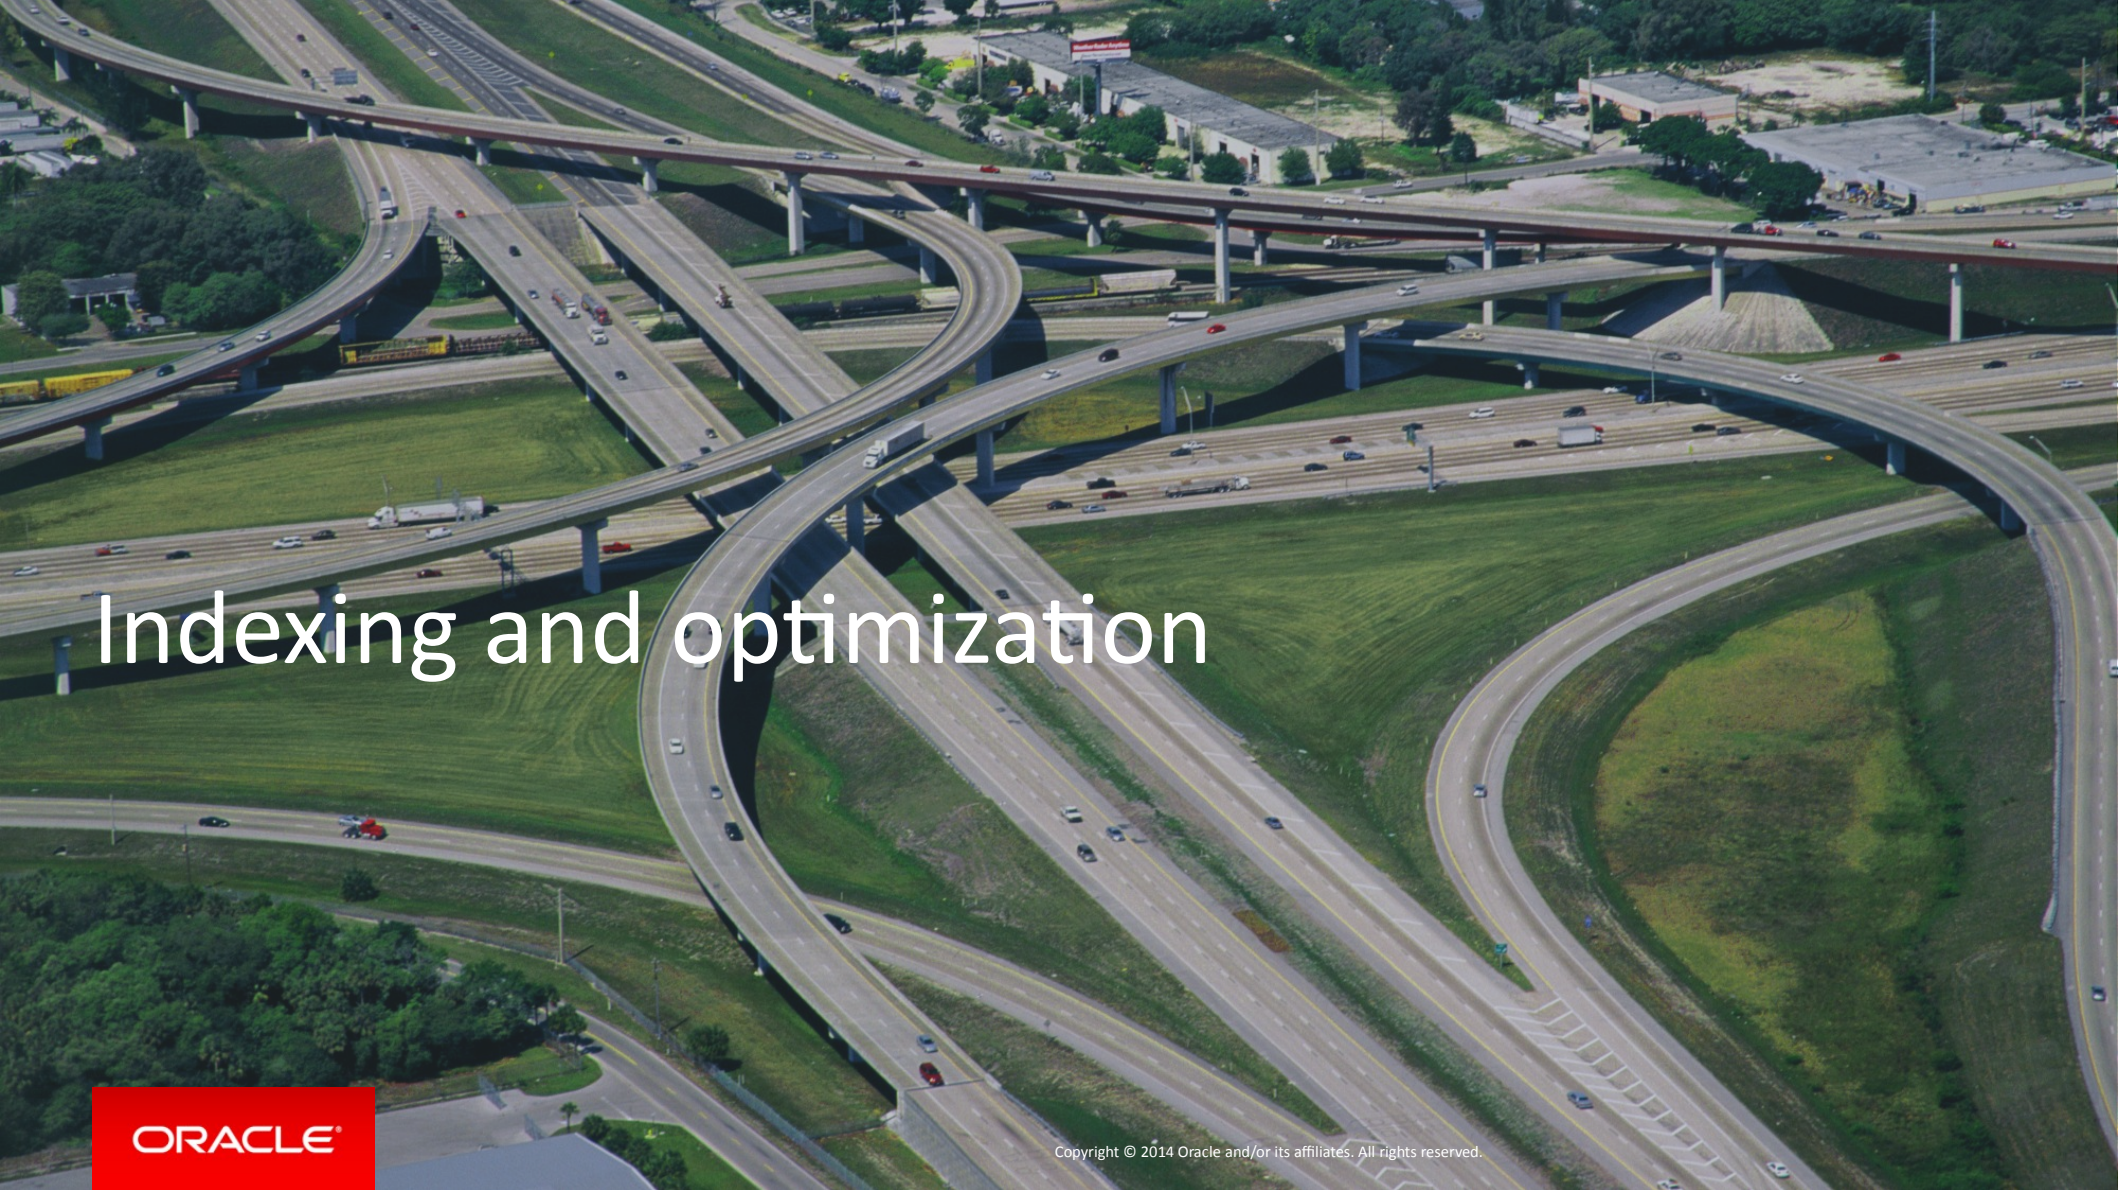
Standing at the Olav Tryggvason Monument, how far is it to the Nidaros Cathedral?

```
SELECT ST_Distance_Sphere(Point(10.3951, 63.4305), Point(10.3958, 63.4269)) AS d;
```

**d**

401.8121469013054





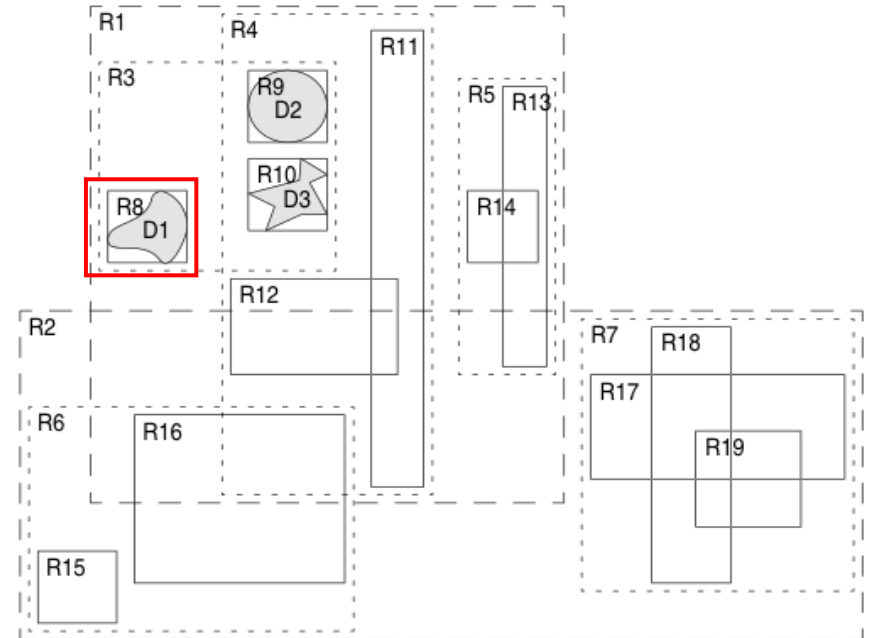
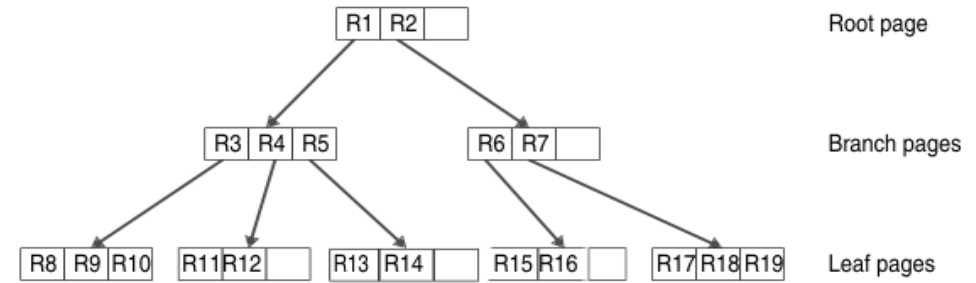
# Indexing and optimization



# Indexes

- R-tree spatial indexes
  - 2d
  - Uses bounding boxes
  - Fully transactional
- Column must be NOT NULL

ALTER TABLE *table* ADD SPATIAL INDEX (*column*);



# Optimization

- The optimizer automatically uses an R-tree index if it thinks it's beneficial
- The query must have a suitable WHERE clause
  - ST\_Contains, ST\_Crosses, ST\_Disjoint, ST\_Equals, ST\_Intersects, ST\_Overlaps, ST\_Touches, ST\_Within
  - MBRContains, MBRDisjoint, MBREquals, MBRIntersects, MBROverlaps, MBRTouches, MBRWithin

# Example

```
SELECT description FROM sights WHERE ST_Within(pos, ST_Buffer(Point(10.3951, 63.4305), 0.004));
```

```
mysql> SHOW STATUS LIKE 'Handler_read%';
```

Variable_name	Value
Handler_read_first	1
Handler_read_key	1
Handler_read_last	0
Handler_read_next	0
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	16

7 rows in set (0.00 sec)

Without index

```
mysql> SHOW STATUS LIKE 'Handler_read%';
```

Variable_name	Value
Handler_read_first	0
Handler_read_key	1
Handler_read_last	0
Handler_read_next	3
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	0

7 rows in set (0.00 sec)

With index

# Example: Displaying data in a web map

# Displaying data in a web map

- A sightseeing map of Trondheim
- Use our sightseeing database from earlier examples
  - Add a few more points to make it more exciting
- Use OpenLayers to display the map
  - OpenStreetMap tiles
  - Add a marker for each place of interest

**L** Any Linux distro will do  
Ubuntu 14.04 LTS

**A** Whichever Apache version came with my OS

**M** MySQL 5.7 DMR from [repo.mysql.com](http://repo.mysql.com) (5.7.6 or newer)  
Sveta Smirnova's JSON UDFs

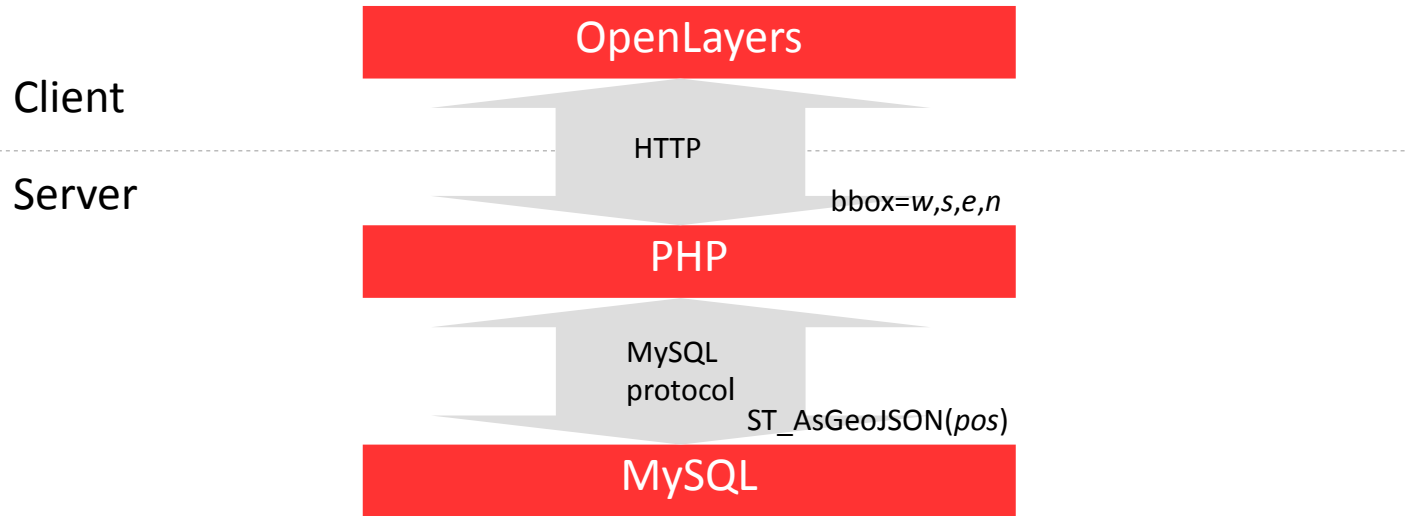
**P** Whichever PHP version came with my OS

# OpenLayers

- Simple setup, following an online guide
  - <http://docs.openlayers.org/library/introduction.html>
- Modifications
  - OpenStreetMap map layer
    - More detailed than the one in the guide
  - Local tile cache (in case I lose my network connection)

# Connecting the parts

- Both MySQL and OpenLayers support GeoJSON
- Use a PHP script to query the database





```
<html><head><title>OpenLayers Example</title>
<script src="openlayers/OpenLayers.js"></script>
</head>
<body>
  <div style="width:100%; height:100%" id="map"></div>
  <script defer="defer" type="text/javascript">
    var map = new OpenLayers.Map('map');
    var osm_url = "http://localhost:8080/map/tiles.php?z=${z}&x={x}&y=${y}&r=mapnik"
    var osm = new OpenLayers.Layer.OSM('osm', [osm_url]);
    map.addLayer(osm);
    map.zoomToMaxExtent();
  </script></body></html>
```

```
<html><head><title>OpenLayers Example</title>
<script src="openlayers/OpenLayers.js"></script>
</head>
<body>
  <div style="width:100%; height:100%" id="map"></div>
  <script defer="defer" type="text/javascript">
    var map = new OpenLayers.Map('map');
    var osm_url =
      "http://localhost:8080/map/tiles.php?z=${z}&x=
        {x}&y=${y}&r=mapnik"
    var osm = new OpenLayers.Layer.OSM('osm',
      [osm_url]);

    var geojson_format = new
      OpenLayers.Format.GeoJSON();
```

```
var sights_layer = new
  OpenLayers.Layer.Vector("Sights",
  {
    strategies: [new OpenLayers.Strategy.BBOX()],
    protocol: new OpenLayers.Protocol.HTTP(
      {
        url: 'sights.php',
        format: geojson_format
      })
  });

map.addLayers([osm, sights_layer]);
map.zoomToMaxExtent();
</script></body></html>
```

# GeoJSON

```
{
  "type": "FeatureCollection",
  "features": [
    {
      "type": "Feature",
      "geometry": {"type": "Point", "coordinates": [10.4017, 63.4282]},
      "properties": {"description": "Old Town Bridge"}
    },
    {
      "type": "Feature",
      "geometry": {"type": "Point", "coordinates": [10.3958, 63.4269]},
      "properties": {"description": "Nidaros Cathedral"}
    }
  ]
}
```

# GeoJSON

One row of the table

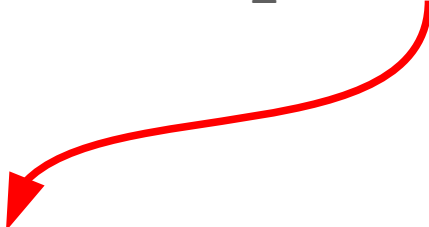
```
{
  "type": "FeatureCollection",
  "features": [
    {
      "type": "Feature",
      "geometry": {"type": "Point", "coordinates": [10.4017, 63.4282]},
      "properties": {"description": "Old Town Bridge"}
    },
    {
      "type": "Feature",
      "geometry": {"type": "Point", "coordinates": [10.3958, 63.4269]},
      "properties": {"description": "Nidaros Cathedral"}
    }
  ]
}
```



# GeoJSON

```
{  
  "type": "FeatureCollection",  
  "features": [  
    {  
      "type": "Feature",  
      "geometry": {"type": "Point", "coordinates": [10.4017, 63.4282]},  
      "properties": {"description": "Old Town Bridge"}  
    },  
    {  
      "type": "Feature",  
      "geometry": {"type": "Point", "coordinates": [10.3958, 63.4269]},  
      "properties": {"description": "Nidaros Cathedral"}  
    }  
  ]  
}
```

ST\_AsGeoJSON(pos)



# JSON\_APPEND

JSON\_APPEND(*object*, *attribute\_name*, *value*)

```
SELECT JSON_APPEND('{}', 'foo', '"bar"') AS json;
```

```
json  
{"foo": "bar"}
```

# Query

```
SELECT JSON_APPEND  
(  
  JSON_APPEND('{"type":"Feature"}', 'geometry', ST_AsGeoJSON(pos)),  
  'properties',  
  JSON_APPEND('{}', 'description', CONCAT(' ', description, ' '))  
) AS json  
FROM sights  
WHERE ST_Within(  
  pos,  
  ST_MakeEnvelope(Point(west, south), Point(east, north))  
);
```

\$bbox = \$\_GET['bbox'];





# The future of MySQL GIS



# General goals

- Provide adequate GIS for existing MySQL users
  - Growing mobile market and emerging IoT market both require spatial features
- Competing with PostGIS in the FOSS GIS DBMS market
  - For basic/common GIS use cases
- Competing with Microsoft SQL Server in the commercial GIS market

# Future enhancements

- A non-flat Earth
  - Ellipsoidal Earth model
  - Projections
- OGC standard metadata tables (SPATIAL\_REF\_SYS, etc.)
- 3d and 4d support
  - 3dm, 3dz, 3dzm
- What else would **you** like to see?
  - Let us know!

“GIS is a form of digital mapping technology. Kind of like Google Earth, but better.”

— Arnold Schwarzenegger, Governor of California

# **Hardware and Software** **Engineered to Work Together**

## Safe Harbor Statement

The preceding is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.

ORACLE®